

5-Minute Presentation Guide

RedM Command Center V3 — Live UI Demo

A timed script for presenting the AI-powered anti-trafficking campaign generator to RedM stakeholders. Includes exact dialogue, screen actions, navigation order, and anticipated Q&A.

Audience	RedM Stakeholders
Format	Live browser demo (design_v3.html)
Duration	5 minutes
Date	April 19, 2026



Before You Start

Setup Checklist

- ☐ Open `design_v3.html` in Chrome (full screen — **F11**)
- ☐ Ensure **Light mode** is selected (better for projectors)
- ☐ If you want a live scan: run backend (`cd backend && python server.py`)
- ☐ If backend is NOT running, the static UI still shows all demo data — skip live scan
- ☐ Test your screen sharing / projector connection
- ☐ Have this guide printed or on a second screen

💡 DELIVERY TIPS

- Speak **slowly and confidently** — 5 minutes is plenty of time
- **Point at the screen** when mentioning specific elements
- Don't read captions word-for-word — **summarize** what they say
- End with the hotline number — it **grounds the mission**
- The **Calendar is the star** — spend the most time there

Navigation Order (Quick Reference)

TIME	TAB	DURATION	WHAT TO SHOW
0:00	Onboarding	30 sec	State the problem + one-sentence pitch
0:30	Intelligence	45 sec	Scan + story cards + "Generate Campaign"
1:15	Calendar ★	75 sec	5-day grid + carousel + captions
2:30	Calendar	45 sec	Approve / Edit / Regen flow
3:15	Rules	45 sec	3-tier compliance system
4:00	Impact → Library	30 sec	KPIs + past campaigns
4:30	—	30 sec	Closing statement + hotline



The Script

0:00 - 0:30

Opening — The Problem

ONBOARDING

Stay on the welcome overlay or Calendar view

*"This is the **RedM Command Center** — an AI-powered tool we built to solve a real problem: anti-trafficking organizations need to post consistent, ethical social media content, but they're understaffed and stretched thin.*

*This system takes a **single news story** and automatically generates a **full week of social media content** — 5 posts, each with captions, AI-generated images, and compliance checks — ready for human review in under 3 minutes."*

SCREEN ACTION

If the onboarding overlay shows, click **"Get Started"** through all 3 steps — they explain the workflow. Then dismiss.

0:30 - 1:15

Intelligence Feed

INTELLIGENCE

Click "Intelligence" in the sidebar

*"Everything starts here — the **Intelligence Feed**. The system scans 12 trusted news sources — Reuters, AP, BBC, Polaris Project, the DOJ — for breaking human trafficking stories.*

*Each story is ranked and displayed as a card. The team picks the most relevant one and clicks '**Generate Campaign**'. That's it — one click triggers the entire AI pipeline."*

POINT OUT

- The **"Scan Now"** button at the top
- The **AI Pick** badge on the top story
- The **"Generate Campaign"** button on any story card
- The **Manual URL** input — "They can also paste any article URL directly"

1:15 - 2:30

Content Calendar ★ THE MAIN EVENT

CALENDAR

Click "Calendar" in the sidebar

*"This is the heart of the system — the **Content Calendar**. From one news story, the AI generates 5 posts for the week, each with a different purpose:"*

Point at each day column as you say:

DAY	TYPE	WHAT TO SAY
Monday	Reactive	"The breaking news hook. Gets immediate attention."
Tuesday	Stat Card	"A key data point from the story — builds credibility."
Wednesday	Explainer	"Educational deep-dive — what does this mean?"
Thursday	Action	"A direct call to action — call the hotline, report suspicious activity."
Friday	Hope	"A positive progress story — prevents compassion fatigue."

*"Every post has an **AI-generated image** tailored to the audience, a **quality score** from our AI judge, and a **compliance badge** confirming it meets ethical standards."*

DEMO ACTIONS (DO THESE IN ORDER)

1. Click on a post image → Carousel Modal opens — "This is the review view"
2. Show the headline, caption, quality score, compliance badge
3. Click → to navigate to the next post
4. Close the carousel (✕ or Esc)
5. Click "Caption" on any post to expand the full text
6. Point out the **character counts** — "IG: /2200, X: /280"

2:30 - 3:15

Approve & Edit Flow

CALENDAR

Stay on Calendar view

*"The system is **human-in-the-loop** — nothing publishes without approval. The reviewer can:*

- **Approve** each post individually
- **Edit** any caption directly
- **Regenerate** an image with custom instructions
- Or **batch-approve all 5** with one click"

DEMO ACTIONS

1. Point at the **"Approve"** button on a post
2. Point at the **"Approve All 5 Posts"** button (top right)
3. Click **"Edit"** on one post → show the edit modal → click **Cancel**

*"Once approved, they can **download the entire campaign as a ZIP** or **translate it** to Spanish, French, or Portuguese — one click."*

POINT AT

The **Download** and **Translate** buttons in the top action bar.

3:15 - 4:00

Content Rules — The Differentiator

RULES

Click "Rules" in the sidebar

*"This is what makes this system different from generic AI tools. We built in **ethical guardrails** — a rules engine that every piece of content must pass through."*

Point at each tier as you explain:

TIER	COLOR	WHAT TO SAY
Legal & Compliance	Red (locked)	"These are locked — can't disable them. Never identify survivors. Always cite sources."
Ethical Language	Orange	"The system automatically replaces harmful terms. For example, it won't say 'rescued' — it says 'identified', because 'rescued' implies the survivor lacked agency."
Brand Voice	Gray	"Customizable rules — 'no sensationalized imagery', 'at least one hope post per week.'"

*"The compliance checker runs **both** a regex-based quick check and an AI deep analysis. If it finds a violation, it **auto-fixes it** before the human even sees it."*

POINT AT

The **98% compliance score** bar at the top of the Rules page.

4:00 - 4:30

Impact & Library (Quick)

IMPACT

Click "Impact" — stay for 10 seconds

"The Impact page tracks reach, engagement, hotline clicks, and API costs per campaign."

LIBRARY

Click "Library" — stay for 15 seconds

"And the Content Library stores every campaign ever generated — with thumbnails. Click any past campaign to reload it into the calendar."

4:30 - 5:00

Closing

*"To summarize: **one news story in, five platform-ready posts out** — each written, illustrated, scored, and compliance-checked by AI, with full human control."*

*The system uses **GPT-4o** for strategic planning, **GPT-4o-mini** for writing and quality, and **Google Gemini** for image generation. Everything is logged, auditable, and translatable.*

*We built this so a team of **2-3 volunteers** can produce what used to take a social media agency a full week."*

Final beat — point at the sidebar:

*"And that number at the bottom — **1-888-373-7888** — that's the National Human Trafficking Hotline. Every single post we generate includes it. That's why we built this."*



With Continued Support

What you've seen today is the **working MVP**. With continued support, here is what the next phase of development unlocks — each item directly multiplies the system's impact and operational maturity.

1. Agent Security & Data Protection

Current state: API keys stored in .env, open CORS, no authentication layer.

ENHANCEMENT	WHAT IT MEANS
LLM Prompt Hardening	Injection guards and output sanitization to prevent adversarial manipulation of AI-generated content — critical for a system producing public-facing material about vulnerable populations.
Role-Based Access Control	Separate Creator, Reviewer, and Admin roles with JWT authentication. Only Admins can modify compliance rules; Reviewers can approve; Creators can generate.
API Key Vault	Migrate from .env files to a secrets manager (AWS Secrets Manager or Azure Key Vault) with key rotation and access logging.
Rate Limiting & Abuse Prevention	Per-user generation limits, IP-based throttling, and anomaly alerts to prevent API cost runaway.
Audit Encryption	End-to-end encryption of the audit trail database. Ensure compliance with data protection regulations for sensitive trafficking intelligence.

2. Sophisticated Content & Higher-Quality Images

Current state: Gemini 3.1 Flash for images, basic Pillow text overlay, 3 font files.

ENHANCEMENT	WHAT IT MEANS
Premium Image Models	Upgrade to DALL-E 3, Midjourney API, or Gemini 2.0 Pro for photorealistic, publication-quality imagery with better composition and lighting.
Brand Template System	Pre-designed templates with RedM logo placement, consistent typography, gradient overlays, and watermarks — replacing the current basic Pillow compositing.
Video Content Generation	Extend from static images to short-form video content (15s Reels/TikToks) using AI video models — the highest-engagement format on modern platforms.
Multi-Story Campaigns	Generate campaigns that weave multiple news stories into a unified narrative arc, rather than one-story-per-week.
A/B Testing Built In	The PromptRegistry already tracks quality scores. Wire it to automatically rotate high-performing prompts and caption styles based on real engagement data.

3. Post-Publish Monitoring & Analytics

Current state: Impact page shows placeholder KPIs; no live social media data.

ENHANCEMENT	WHAT IT MEANS
Live Social Media Integration	Connect Instagram Graph API, X API v2, LinkedIn Marketing API, and Facebook Pages API to pull real-time impressions, likes, shares, saves, and comments.
Hotline Click Attribution	UTM-tagged links in every post. Track exactly how many hotline visits and calls were generated by each campaign — the most important metric.
Sentiment Analysis Dashboard	AI-powered comment monitoring. Detect negative reactions, misinformation in replies, or requests for help — with automated escalation alerts.
Content Performance Feedback Loop	Feed engagement data back into the Campaign Planner so the AI learns which angles, tones, and content types perform best with each audience profile.
Weekly Report Generation	Automated PDF performance reports for stakeholders — reach, engagement, cost-per-impression, and hotline attribution — sent every Monday.

4. Operational Cost Optimization

Current state: ~\$0.05-0.15 per campaign (43-48 LLM calls). Gemini images are on free tier.

ENHANCEMENT	IMPACT
Intelligent Caching	Cache common LLM responses (audience matching, compliance checks for similar content). Could reduce LLM calls by 30-40% for recurring topics.
Model Routing Optimization	Move simple tasks (hashtag generation, character counting) to local models or rule-based logic. Reserve GPT-4o for planning only.
Batch API Calls	Use OpenAI's Batch API (50% cost reduction) for non-urgent tasks like quality scoring and compliance checking.
Cost Budget Alerts	Set monthly spending caps with email/Slack alerts at 50%, 80%, and 100% thresholds.
Monthly Cost Projection	At current rates: 4 campaigns/month = ~\$0.40-0.60/month in AI costs. At scale (20 campaigns/month) ≈ \$2-3/month. Infrastructure (hosting) adds ~\$5-10/month for a basic cloud deployment.

5. Direct Social Media Publishing

Current state: ZIP download only. No API-level publishing.

ENHANCEMENT	WHAT IT MEANS
One-Click Publish	After human approval, publish directly to Instagram, X, LinkedIn, and Facebook via their official APIs — no manual copy-paste.
Scheduled Publishing	The APScheduler infrastructure already exists. Wire it to platform APIs for time-optimized posting (best engagement windows per platform).
Cross-Platform Adaptation	Automatically adjust caption length, hashtag count, and image format per platform's algorithm preferences.

 ESTIMATED TIMELINE WITH SUPPORT

PHASE	DURATION	DELIVERABLES
Phase 1	2-3 weeks	Security layer (auth, RBAC, key vault) + font fix + direct publishing MVP
Phase 2	3-4 weeks	Premium images + brand templates + monitoring dashboard + live API connections
Phase 3	4-6 weeks	Video content + feedback loop + cost optimization + weekly reports



If They Ask Questions

QUESTION	YOUR ANSWER
"How long does generation take?"	"About 2-3 minutes for all 5 posts with images."
"What if the AI writes something wrong?"	"Every post passes through a Quality Gate (scores 0-10) and a Compliance Checker that auto-fixes banned terms. Nothing publishes without human approval."
"Can we customize the rules?"	"Yes — the Brand Voice tier is fully editable. Ethical language rules are toggleable. Legal rules are locked by design."
"What AI models does it use?"	"GPT-4o for campaign planning, GPT-4o-mini for writing and quality checks, and Google Gemini for image generation."
"How much does it cost per campaign?"	"Roughly \$0.05–0.15 per campaign. At scale, 20 campaigns/month costs about \$2–3 in AI. Hosting adds ~\$5-10/month."
"Can it post directly to social media?"	"Not yet — currently ZIP downloads. Direct API publishing is the first item in Phase 1 of our roadmap."
"What languages does it support?"	"English, Spanish, French, and Portuguese. One-click translation from the calendar view."
"What news sources does it scan?"	"12 trusted domains: Reuters, AP, BBC, The Guardian, Al Jazeera, UNODC, Polaris Project, Human Rights Watch, IJM, ILO, and Thorn."
"Is the system secure?"	"The MVP runs locally. Our Phase 1 roadmap adds JWT auth, role-based access control, and a secrets vault before any cloud deployment."
"How do we measure real impact?"	"Phase 2 adds live social API connections and UTM-tagged hotline links — we'll track exactly how many calls each campaign generates."

Remember

 **Start Strong**
Lead with the problem, not the tech

 **Calendar is the Star**
Spend the most time there (75 sec)

 **Emphasize Ethics**
The Rules page is your differentiator

 **Close with the Roadmap**
Show what's possible with support

Good luck! 

1-888-373-7888 | Text BeFree to 233733